



OPEN Knowledge graph driven medicine recommendation system using graph neural networks on longitudinal medical records

Rajat Mishra¹ & S. Shridevi²✉

Medicine recommendation systems are designed to aid healthcare professionals by analysing a patient's admission data to recommend safe and effective medications. These systems are categorised into two types: instance-based and longitudinal-based. Instance-based models only consider the current admission, while longitudinal models consider the patient's medical history. Electronic Health Records are used to incorporate medical history into longitudinal models. This project proposes a novel Knowledge Graph-Driven Medicine Recommendation System using Graph Neural Networks, KGDNNet, that utilises longitudinal EHR data along with ontologies and Drug-Drug Interaction knowledge to construct admission-wise clinical and medicine Knowledge Graphs for every patient. Recurrent Neural Networks are employed to model a patient's historical data, and Graph Neural Networks are used to learn embeddings from the Knowledge Graphs. A Transformer-based Attention mechanism is then used to generate medication recommendations for the patient, considering their current clinical state, medication history, and joint medical records. The model is evaluated on the MIMIC-IV EHR data and outperforms existing methods in terms of precision, recall, F1 score, Jaccard score, and Drug-Drug Interaction control. An ablation study on our models various inputs and components to provide evidence for the importance of each component in providing the best performance. Case study is also performed to demonstrate the real-world effectiveness of KGDNNet.

Keywords Medicine recommendation, Graph neural network, Knowledge graphs, Attention mechanism

Medicine recommendation has been an important area of research in the last few years. The purpose of medicine recommendation systems is to assist healthcare professionals to analyse a patient's admission data regarding diagnoses, illnesses, medical procedures to prescribe a set safe and accurate medications that will help in mitigating the patient's illness. These systems prove to be especially useful when patients are diagnosed with multiple illnesses and undergo several procedures during their stay. However, as doctors and researchers discover new diseases and develop new medicines and procedures, they have to take a number of factors into account while recommending medicines to a patient.

Firstly, they have to consider the patient's medical history while prescribing medication for the current admission. Medicine recommendation models can be divided into two categories, *instance-based* and *longitudinal-based*. *Instance-based* models provide medicine recommendations by taking only the current admission record into account^{1,2}. These models do not consider the patient's medical history and hence cannot take into consideration any previous diagnoses and medications that could affect the recommendations. This affects the accuracy and effectiveness of such models. Recently, to address these issues, longitudinal-based recommendation methods such as^{3–7} have been proposed, which leverage the temporal dependencies present in the patient's medical history to provide personalised and safer recommendations. Electronic Health Records (EHRs) are used to incorporate medical history into *longitudinal-based* models. EHRs like MIMIC^{8,9} systematically collect historical medical information of a patient including diagnoses, procedures, prescription among others in the form of medical codes across admissions^{2,10}.

Secondly, they have to take into consideration the presence of Drug-Drug Interactions (DDIs)¹¹ between different pairs of medicines. According to National Institutes of Health, DDIs can “change a drug's effect on the body when the drug is taken together with a second drug”. These pairs may interact affect the action of either or both drugs and lead to adverse reactions which can deteriorate a patient's health condition^{12,13}. For example, a

¹School of Computer Science and Engineering, Vellore Institute of Technology - Chennai, Chennai, India. ²Centre for Advanced Data Science, Vellore Institute of Technology - Chennai, Chennai, India. ✉email: shridevi.s@vit.ac.in

patient suffering from Adenovirus Pneumonia can be prescribed Ertapenem and another patient suffering from Coronavirus Pneumonia can be prescribed Linezolid. However, if a patient is prescribed both medicines, the DDI between them can result in edema/sepsis in the patient.

Various deep learning approaches capture diverse relationships between medical entities and patients, encompassing factors such as DDIs, ontologies and semantic information from various external biomedical sources, leading to creation of Medical Knowledge Graphs (KGs). Medical KGs consist of various types of nodes representing diagnoses, procedures, medicines and patients with multiple relations connecting them together. LEAP¹ introduces an *instance-based* deep learning framework that, given a patient's current medical information, utilises an attention mechanism to prescribe a safe combination of medicines by taking external DDI information into consideration. RETAIN¹⁴ proposes a “reverse time attention” system using two layers of RNN to take past admissions into account while creating a recommendation set for current admission. GAMENet³ takes into account various drug-drug interactions by adopting a dynamic “memory module” implemented using Graph Convolution Networks (GCN) which learns EHR and DDI KGs and uses RNNs to learn patient history. SafeDrug⁴ leverages MIMIC-III, DrugBank and molecular structure of drugs by utilising a bipartite learning module to provide safe drug recommendation with a significant reduction in DDIs. MICRON⁵ uses a “recurrent residual network” that updates a patient's medical information and propagates them to the next visit to preserve the patient's temporal information. COGNet⁶ introduces a “copy-or-predict mechanism” to that uses the patient's medical history along with current diagnoses to determine whether to prescribe (“copy”) a previous medicine or recommend a new medicine. COGNet also implements a Transformer Encoder mechanism¹⁵, which consists of Self-Attention followed by Layer Normalisation¹⁶ to learn the patient's temporal medical history.

4SDrug¹⁷ uses a set-oriented representations, similarity metrics, and importance-based set aggregation for symptom analysis, along with intersection-based set augmentation for recommendation. DrugRec¹⁸ addresses recommendation bias by modelling causal inference to address bias, track patients' health history, and manage DDIs by solving them as a satisfiability problem. MoleRec¹⁹ employs a hierarchical architecture to model molecular inter-substructure interactions and their effect on a patient's health condition. ACDNet⁷ uses the attention-based Transformer Encoder along with GAT²⁰ and GCN²¹ to learn the medication KGs and uses cosine similarity along with Feed-Forward Neural Networks for recommending medicines. PROMISE²² encodes EHR hypergraphs to learn patient representations and acquires patient representations with semantic information by encoding clinical texts from EHR. VITA²³ introduces two novel ideas: relevant visit selection that excludes past visits deemed irrelevant and target aware attention to capture the relevancy between current and past visits.

Using the above findings, we propose a novel Knowledge Graph-Driven Medicine Recommendation System using Graph Neural Networks on Longitudinal Medical Records (KGDNet), as illustrated in Fig. 5a., that utilises the EHR data along with ontologies and DDI knowledge from external sources to construct admission-wise clinical and medicine KGs for every patient using a variety of relations including semantic, ontological and DDI relations. These pairs of KGs for each admission will be utilised to generate temporal information to be used for providing personalised recommendation for the patient. We also create a drug knowledge base by converting the DDI information into a KG which will be leveraged to provide safe recommendations. These KGs are then learned using GNNs designed for highly relational data to generate embeddings. We then exploit RNNs to learn temporal data from medical record embeddings from each admission and generate temporal features from clinical and medicine streams across admissions. These temporal features are then combined using a Fusion Module utilising Multi-Layer Perceptrons (MLPs) to produced joint temporal features for each admission. These features are then passed onto the Recommendation module that uses a Transformer-based Attention mechanism that uses the joint temporal features, the clinical features from previous admissions and the current clinical state of a patient to prescribe a set of medicines. We utilised the EHR cohort MIMIC-IV⁹ along with DDI data¹¹ and medical ontologies ICD and ATC from external knowledge bases to evaluate our model. We compare the performance of our model against various of deep learning medicine recommendation models. The results show that our model consistently outperforms its counterparts across multiple performance metrics including Precision-Recall Area Under Curve, Jaccard and F1 scores, and proves to be excellent at reducing DDIs while also maintaining consistently high performance.

Our main contributions are summarised as follows:

1. We propose a novel KGDNet framework for safe and effective medicine recommendation that maximises accuracy and minimises Drug-Drug Interactions within the medicine sets by exploiting semantic, relational and ontological knowledge to construct admission-wise medical KGs for each patient in the EHR. We use GNNs to learn the patient's medical data through the KGs. Global DDI KG is learned and removed from patient's medication embeddings.
2. In order to learn a patient's admission history, we utilise Recurrent Neural Networks to hierarchically learn the clinical and medicine streams using KG embeddings obtained from the GNNs and then fuse the temporal features of each stream to construct a joint medical stream using Multi-Layer Perceptrons.
3. Using a Transformer-based Attention mechanism, we create a recommender module that uses a patient's current clinical state and medication history along with joint medical records and clinical history to generate medication recommendations for the patient.
4. We optimise our model using a combined loss method that takes multi-label prediction accuracy and DDI rate into account. We evaluate our model on the MIMIC-IV EHR data using various performance metrics to demonstrate the effectiveness of our proposed KGDNet in comparison to existing methods.

Results
Model prediction

To evaluate the performance of our novel medication recommendation model KGDNet, we compare our method with various baselines, divided into two groups, instance-based, including Logistic Regression (LR), LEAP¹, and longitudinal-based, including RETAIN¹⁴, GAMENet³, SafeDrug⁴, MICRON⁵, COGNet⁶, DrugRec¹⁸, MoleRec¹⁹, ACDNet⁷, PROMISE²² and VITA²³. Further information is provided in Section 4.6.

We compare predictive performance of the methods on the MIMIC-IV⁹ EHR cohort, after performing data processing as mentioned in Section 4.1, on the basis of PR-AUC, Jaccard and F1 metrics. We also test the safety of our model on the basis of DDI rates among the recommended medicines. We report KGDNet’s performance comparison against baselines in Table 1 and illustrate our results in Fig. 1. We set our DDI threshold to 0.08 to reflect the real-world dataset. *Instance-based* methods, LR (PR-AUC=0.7090) and LEAP (PR-AUC=0.5506), performed poorly in comparison to the *longitudinal-based* models as they do not take patient history into consideration.

The longitudinal-based models perform better as they take into consideration the patient’s historical information along with the current clinical state using a variety of methods. While RETAIN (PR-AUC=0.7154, DDI=0.0904) only uses longitudinal information, GAMENet (PR-AUC=0.7487, DDI=0.0848) achieves better results by introducing DDI and drug co-occurrence information using GNNs and SafeDrug (PR-AUC=0.7503, DDI=0.0737) incorporates molecular structures of drugs and succeeds in attaining higher scores and reducing DDI. MICRON (PR-AUC=0.6842) introduced a “recurrent residual network” to tackle redundancy but fails to account for co-occurrence relationship among medicines from previous visits. COGNet (PR-AUC=0.7525) uses a “copy-or-predict” system alongside Transformers¹⁵ to consider the relationship between current and previous prescription records. 4SDrug (PR-AUC=0.7011, DDI=0.0637) achieves a good DDI Rate by utilises a set-based similarity measurement to generate small, concise medicine sets but sacrifices accuracy in doing so. DrugRec (PR-AUC=0.7225, DDI=0.0633) uses casual inference to address recommendation bias and track patient history get low DDI rates and good accuracy. MoleRec (PR-AUC=0.7288, DDI=0.0695) models the relation between the molecular structures and patient’s clinical information to achieve good accuracy. ACDNet (PR-AUC=0.7501, DDI=0.0849) uses Transformers, to learn the diagnoses, procedures and medication history of the patient separately and then fuses the them using Transformers too, similar to COGNet. PROMISE (PR-AUC=0.7335, DDI=0.0621) creates hypergraphs from EHR and encodes semantic information from clinical texts to achieve very good balance between DDIs and accuracy. VITA (PR-AUC=0.7225, DDI=0.0922) captures the relevancy between past and current visits to get good accuracy but high DDI rates.

However, KGDNet outperforms ACDNet by using Multi-Head Attention (MHA)¹⁵, similar to Transformers, instead of Cosine Similarity to generate recommendations and using GRU along with MHA to learn the patient’s history from KG embeddings. Unlike the baselines, to account for the highly relational clinical data, we create a clinical KG along with the DDI and medicine KGs for each visit, which are learned using GNNs and passed to RNNs. We also use an attention mechanism for prescribing medicines unlike previous works. This results in superior scores in terms of both predictive efficiency and safety for KGDNet (PR-AUC=0.7657, DDI=0.0665).

During training, as shown in Fig. 2, the PRAUC, F1, and Jaccard scores all steadily increased with the number of epochs. The DDI rate, however, exhibited a more complex trend. Initially, it decreased as the number of epochs increased. However, it then began to increase before relatively stabilising at a later point in the training process. This trend along with the variations every few epochs in the DDI rate in Fig. 2d. is likely due to the fact that different sets of medicines are recommended during each epoch. The eventual stabilisation of the DDI rate suggests that the recommendation sets are becoming more consistent as training progresses.

Model	DDI Rate	PRAUC	F1 Score	Jaccard	Avg. # of meds
LR	0.0762 ± 0.0004	0.7090 ± 0.0014	0.6007 ± 0.0013	0.4510 ± 0.0013	8.9866 ± 0.0374
LEAP ¹	0.0731 ± 0.0004	0.5506 ± 0.0015	0.5820 ± 0.0012	0.4287 ± 0.0012	11.5198 ± 0.0459
RETAIN ¹⁴	0.0904 ± 0.0011	0.7154 ± 0.0018	0.6170 ± 0.0023	0.4613 ± 0.0026	12.8949 ± 0.0923
GAMENet ³	0.0848 ± 0.0005	0.7487 ± 0.0015	0.6449 ± 0.0017	0.4920 ± 0.0018	19.3289 ± 0.0912
SafeDrug ⁴	0.0737 ± 0.0007	0.7503 ± 0.0013	0.6578 ± 0.0019	0.5065 ± 0.0020	15.9642 ± 0.0335
MICRON ⁵	0.0681 ± 0.0016	0.7124 ± 0.0025	0.6465 ± 0.0032	0.4754 ± 0.0026	15.6963 ± 0.2875
COGNet ⁶	0.0894 ± 0.0003	0.7525 ± 0.0008	0.6467 ± 0.0009	0.4884 ± 0.0009	19.7235 ± 0.0242
4SDrug ¹⁷	0.0637 ± 0.0004	0.7011 ± 0.0011	0.6034 ± 0.0010	0.4539 ± 0.0011	12.5213 ± 0.0665
DrugRec ¹⁸	0.0633 ± 0.0012	0.7225 ± 0.0010	0.6455 ± 0.0007	0.4904 ± 0.0011	15.7565 ± 0.1223
MoleRec ¹⁹	0.0695 ± 0.0012	0.7288 ± 0.0023	0.6452 ± 0.0012	0.5001 ± 0.0015	18.5714 ± 0.1244
ACDNet ⁷¹	0.0849 ± 0.0005	0.7501 ± 0.0017	0.6564 ± 0.0013	0.5077 ± 0.0015	12.7024 ± 0.0005
PROMISE ²²	0.0621 ± 0.0007	0.7335 ± 0.0010	0.6517 ± 0.0008	0.4973 ± 0.0010	17.1309 ± 0.0741
VITA ²³	0.0922 ± 0.0034	0.7225 ± 0.0010	0.6583 ± 0.0007	0.5174 ± 0.0011	14.5454 ± 0.1001
KGDNet	0.0665 ± 0.0010	0.7657 ± 0.0015	0.6765 ± 0.0017	0.5218 ± 0.0018	19.2273 ± 0.0912

Table 1. Performance comparison of Model against baselines on MIMIC-IV EHR dataset. The base DDI rate in EHR test data is 0.0781. ¹ACDNet metrics taken from the paper⁷ itself as their code was not publicly available for experimentation at the time of submission

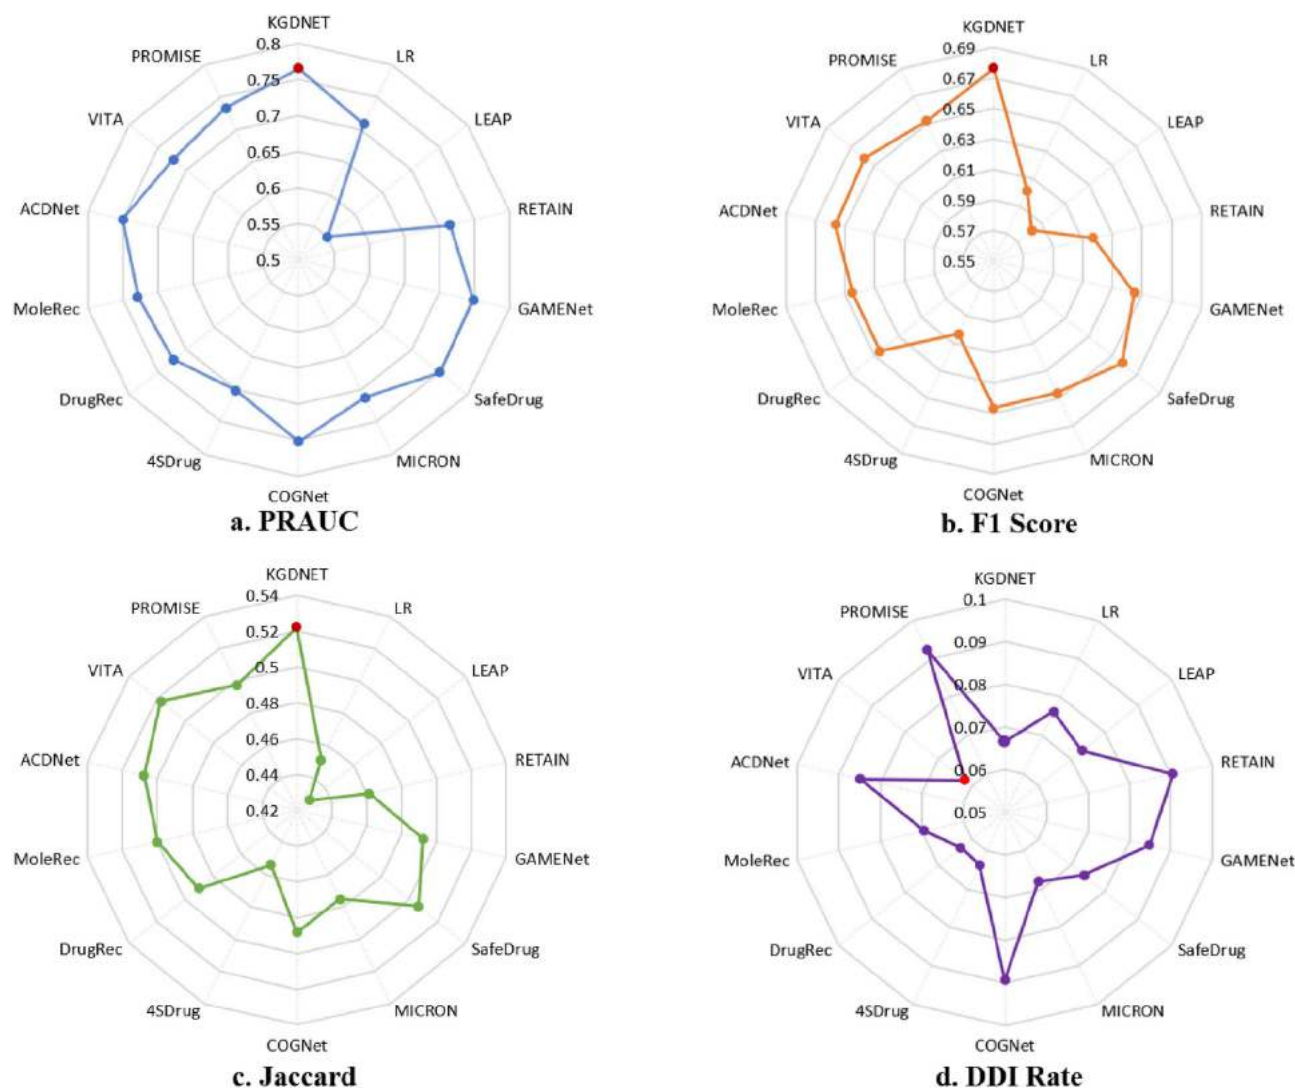


Fig. 1. Evaluation performance of KGDNet against baselines. All results were obtained after 10 rounds of bootstrap sampling with 80% samples. The boxplots compare our model, KGDNet, alongside various baselines on the basis of (a) PRAUC, (b) F1 Score, (c) Jaccard, (d) DDI Rate. The red dot (●) signifies the best result from the models for the respective metric.

Analysis of DDI rate thresholds

Ensuring the safety of medicine recommendations is our primary concern. Hence, we test our model's capability to control DDI rates and show that DDI rates can be effectively controlled by the hyperparameter λ . The DDI rate in the MIMIC-IV EHR cohort is 0.0781. The hyperparameter λ allows us to alter the training loss for each patient such that if the patient-level DDI is less than λ , we only need to maximise the prediction accuracy but if it is greater, then we adjust our loss function to focus on minimising DDIs as well. We test our model with DDI thresholds in the range of [0.05, 0.1] with 0.01 increments. We show that our model is capable of controlling DDI rates. We report KGDNet's performance against DDI threshold in Table 2 and illustrate our results in Fig. 3. In lower DDI thresholds, the model controls the DDI rates well but struggles with achieving good scores due to the low threshold restricting the recommendation size. However, when the threshold is in the range [0.07, 0.1], the model does well to suppress the DDI rates well below the thresholds and performs better. Due to threshold being increased, more medicines are allowed in the recommendation set. Hence, this shows that KGDNet can successfully mimic clinicians when prescribing medicines by balancing the trade-off between DDI rates and accuracy.

Ablation study for feature importance

To perform ablation study, we observe the impact of various modules by removing them from our model and training it.

- **KGDNet w/o K^m :** We remove the medicine knowledge graph from the input and subsequently, we remove the DDI graph.

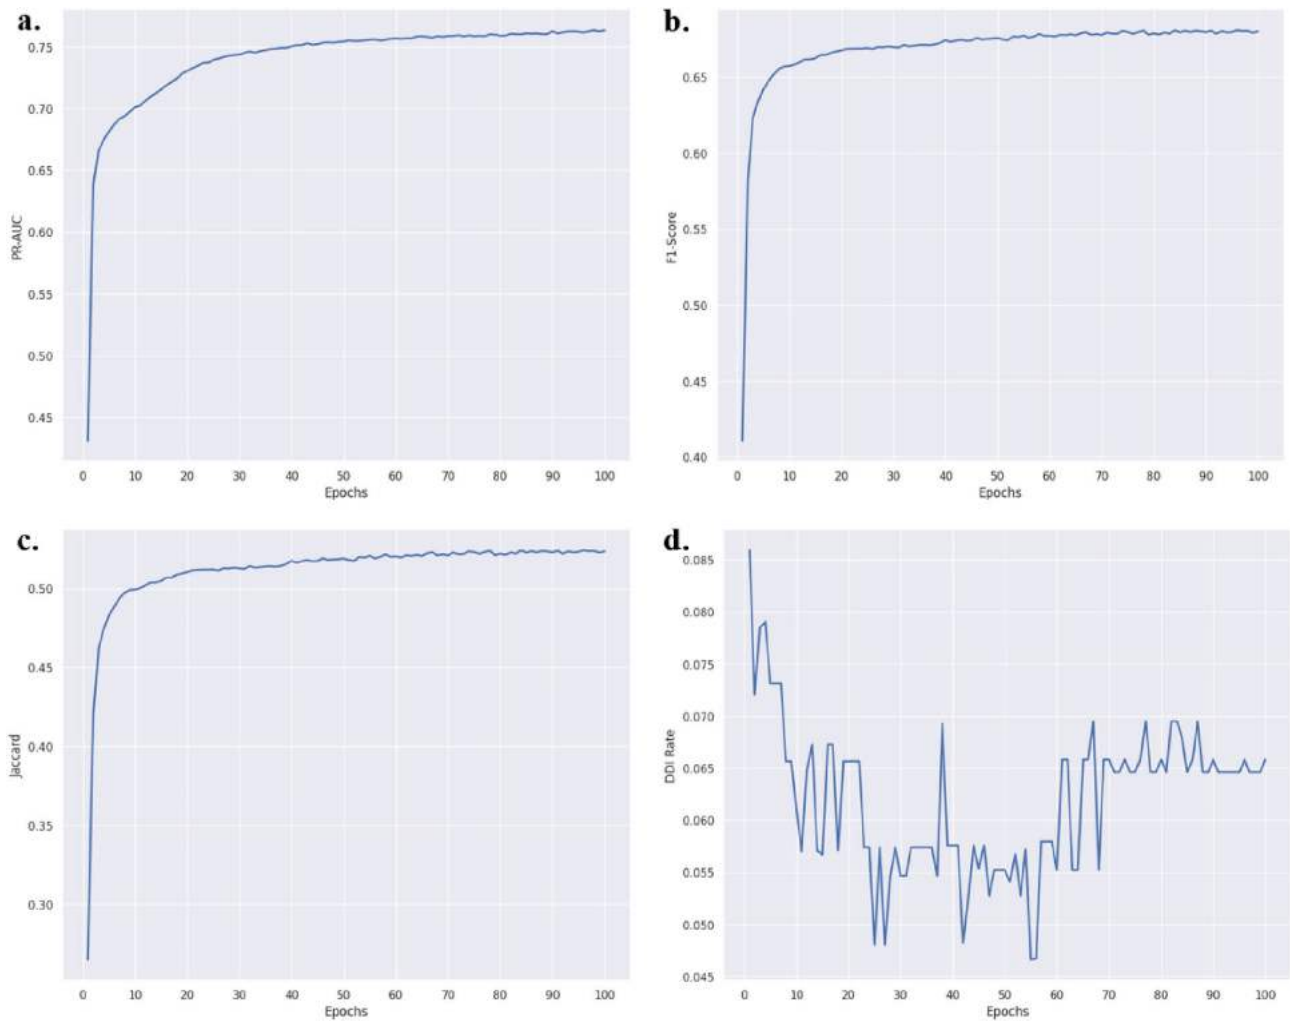


Fig. 2. Values of various metrics during training over the course of 100 epochs: **(a)** PRAUC, **(b)** F1 Score, **(c)** Jaccard, **(d)** DDI Rate.

Threshold	DDI Rate	PRAUC	F1 Score	Jaccard	Avg. # of meds
0.05	0.0527 ± 0.0004	0.7422 ± 0.0023	0.6592 ± 0.0014	0.5095 ± 0.0020	17.0847 ± 0.0663
0.06	0.0618 ± 0.0010	0.7589 ± 0.0015	0.6688 ± 0.0003	0.5136 ± 0.0011	19.0756 ± 0.0735
0.07	0.0658 ± 0.0011	0.7643 ± 0.0018	0.6668 ± 0.0023	0.5166 ± 0.0016	19.1234 ± 0.0923
0.08	0.0665 ± 0.0007	0.7657 ± 0.0015	0.6765 ± 0.0017	0.5218 ± 0.0018	19.2273 ± 0.0912
0.09	0.0703 ± 0.0004	0.7631 ± 0.0013	0.6735 ± 0.0019	0.5191 ± 0.0020	20.2971 ± 0.0335
0.10	0.0734 ± 0.0003	0.7561 ± 0.0008	0.6698 ± 0.0009	0.5149 ± 0.0009	21.2235 ± 0.0242

Table 2. Performance of KGDNet under a spectrum of DDI thresholds.

- **KGDNet w/oDDI:** We remove the DDI knowledge graph and DDI adjacency matrix from the input and subsequently, the DDI loss.
- **KGDNet w/oFusion:** We remove the fusion module for combining the medicine and clinical streams and simply concatenate the two streams.
- **KGDNet w/oAttn.:** We remove the attention module for recommendation and replace it with a mean operation.
- **KGDNet w/oK^m, Attn.:** We remove the medicine knowledge graph from the input and subsequently, we remove the DDI graph and DDI loss. We also remove the attention module for recommendation and replace it with a mean operation. We report the results of our ablation study in Table 3. **Model w/oK^m** discards the medication KG K_{t-1}^m for each admission X_t , subsequently discarding the DDI KG K^{ddi} as it is directly connected to the medication KG. We also discard the related modules GNN^m , GNN^{ddi} and RNN^m and

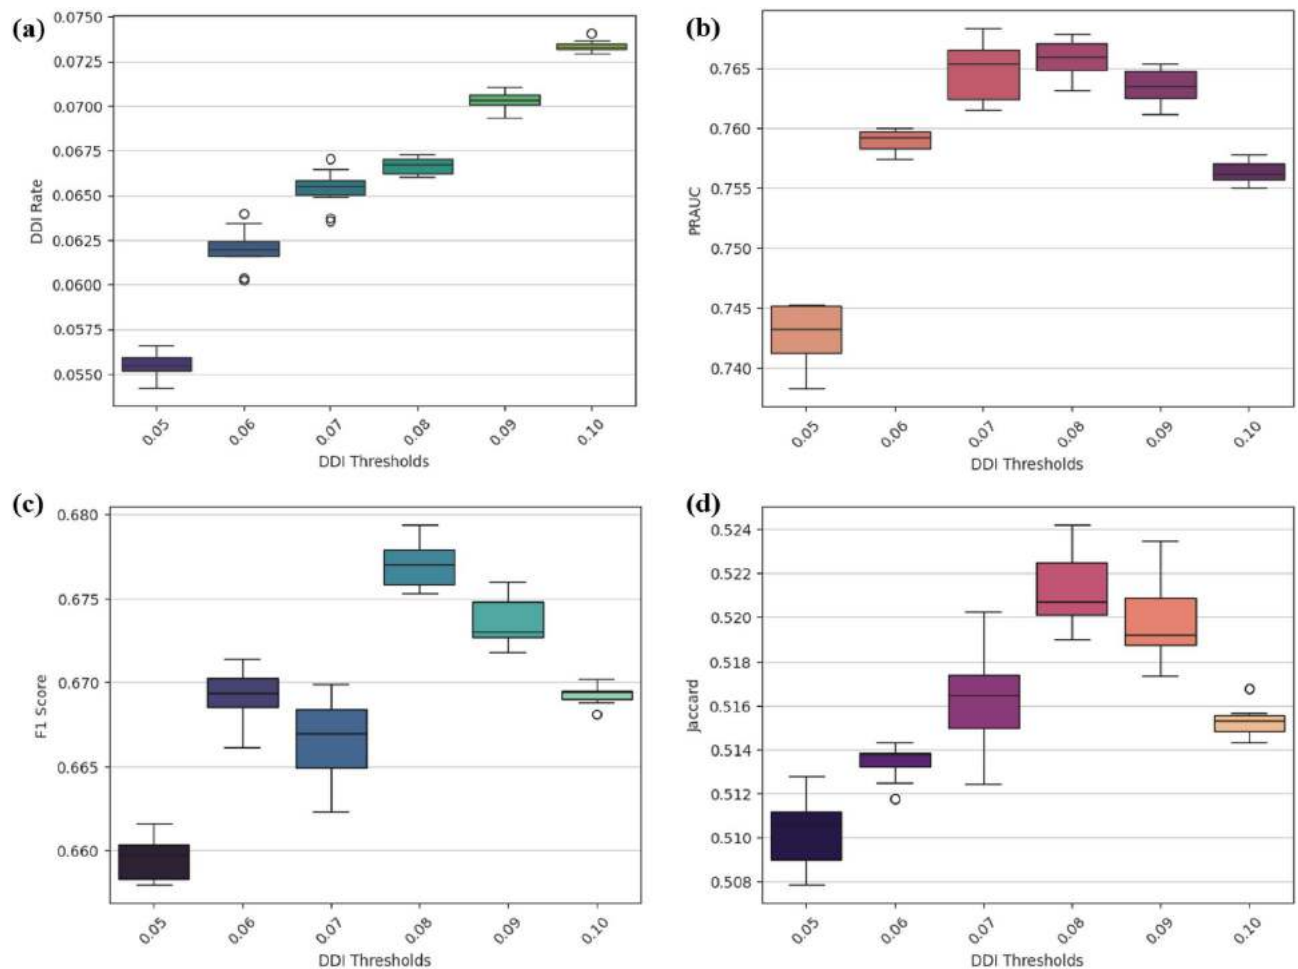


Fig. 3. Evaluation performance of KGDNet against a range of DDI thresholds. The boxplots illustrate the performance of KGDNet on the basis of (a) DDI Rate, (b) PRAUC, (c) F1 Score, (d) Jaccard. We see that DDI Rate is well under the threshold in most cases while rest of the metrics have the highest value when DDI threshold = 0.08.

Model	DDI Rate	PRAUC	F1 Score	Jaccard	Avg. # of meds
KGDNet w/o K^m	0.7538 ± 0.0015	0.6748 ± 0.0013	0.5183 ± 0.0008	0.0679 ± 0.005	19.2692 ± 0.0230
KGDNet w/o DDI	0.7533 ± 0.0018	0.6668 ± 0.0018	0.5174 ± 0.0013	0.0819 ± 0.0028	20.3558 ± 0.0425
KGDNet w/o Fusion	0.7584 ± 0.0015	0.6729 ± 0.0021	0.5207 ± 0.0013	0.0688 ± 0.0005	19.3876 ± 0.0098
KGDNet w/o Attn.	0.7574 ± 0.0023	0.6741 ± 0.0023	0.5199 ± 0.0015	0.0664 ± 0.0014	19.4639 ± 0.0154
KGDNet w/o K^m , Attn.	0.7558 ± 0.0009	0.6629 ± 0.0016	0.5163 ± 0.0010	0.0707 ± 0.0009	18.3763 ± 0.0224
KGDNet	0.7657 ± 0.0015	0.6765 ± 0.0017	0.5218 ± 0.0018	0.0658 ± 0.0011	19.2273 ± 0.0912

Table 3. Ablation Study for Various Components of KGDNet on MIMIC-IV.

the fusion modules. We only utilise the clinical knowledge graph K^c and the attention module uses only the obtained longitudinal clinical features h_t^c in the MHA module. These changes result in lower scores and high DDI rate as the model lacks any knowledge of the medication history and DDI information. In **Model w/o DDI**, we provide patient medical history for each admission but not the DDI information by removing the DDI loss L_{ddi} . This results in slightly higher prediction scores compared to **Model w/o K^m** but due to absence of any DDI information and no feedback from the DDI loss function, this variant suffers from high levels of DDIs in the medication set.

Model w/o Fusion discards the fusion module that takes the longitudinal clinical and medication streams obtained from RNNs and fuses them to create the joint medical stream for a patient. Instead we perform simple concatenation of the clinical and medication streams, i.e., $f_t^{c+m} = h_t^c \cdot h_t^m$, similar to GAMENet³. This

produces slightly lower scores for PRAUC than KGDNet but the other metrics are very similar, showing better performance than other variants. These results verify that incorporating both clinical and medication data are necessary for accurate recommendations. Next, **Model w/oAttn.** discards the MHA module and replaces it with a mean operation, i.e., $l_t = \text{LayerNorm}(e_t^c + \text{mean}(h_t^c, h_{t-1}^{c+m}))$. This produces inferior scores than KGDNet and indicates the importance of MHA in assisting the model focus on key features of the patient's medical state. We further extend this variant by introducing **Model w/oK^m, Attn.** that discards the medical knowledge graph and thus the GNN^m , GNN^{di} and RNN^m and the fusion modules. The attention module is a layer normalisation of the current clinical state along with the longitudinal clinical features of the patient, i.e., $l_t = \text{LayerNorm}(e_t^c + h_t^c)$. This leads to further reduction in scores underlining the importance of the patient's medication history and the need for an Attention module. The performance of the variants are also captured in Fig. 4 by comparing the Precision-Recall and ROC-AUC curves of the variants. We can see from the Precision-Recall curves that the KGDNet model achieves better precision than all ablated variants while the True Positive Rate gradually outperforms the ablated variants at higher rates in the ROC curves.

Discussion

Until recently, studies focused on genetics-based predictions²⁴ to provide personalised recommendations. However, research focus has shifted to recommendations using medical data from Electronic Health Records (EHRs). EHRs collect a variety of medical information relating to diagnoses, procedures and prescriptions from a large number of patients across their admissions. EHRs have become a valuable source of information for deep-learning models^{25,26}. Given their inherent time-series format, EHRs offer personalized insights for each patient regarding various medical entities such as medicine, diagnoses, etc. which can be leveraged to provide medication personalised to the patient's conditions^{27,28}. Reducing Drug-Drug Interactions (DDIs) among medicines has also been an active area of research. DDIs can result in adverse side effects that can deteriorate a patient's health. Various methods have been proposed to control DDIs in medicines. Several Graph Neural Network-methods to mitigate DDIs have been proposed^{29,30} that leverage extensive biomedical networks and knowledge graphs such as DrugBank³¹ and TWOSIDES¹¹.

Graph neural networks for encoding patient EHRs

Graph Neural Networks (GNNs) are deep-learning methods designed to work with heterogeneous graph-structured data^{32,33}. Unlike Convolutional Neural Networks (CNNs), GNNs analyse complex relational data between network entities, capturing graph dependencies through “message passing” between nodes. They are widely used in recommendation models^{34,35}. Models like Graph Convolutional Networks (GCNs)²¹, Graph Attention Networks (GAT)²⁰, and Message Passing Neural Networks (MPNN)³⁶ have been applied by^{3,6,37,38} and⁴, respectively. GAMENet, for instance, uses GCNs to learn medicine knowledge graphs (KGs) with drug-drug interaction (DDI) information and co-occurrences. Notably, no work has applied GNNs and KGs to model *patient* diagnoses and procedure data, which is highly relational. Clinical data from EHRs contain extensive ontological and semantic relationships, providing insights like disease and diagnoses co-occurrences, enriching patient-specific information. Inspired by the need for a way to account for the clinical relational data, we propose personalized medical KGs for both medication and clinical data for each patient admission. To capture these relations, we employ relation-aware GNNs (RAGNNs) to learn medical representations and use RNNs for longitudinal learning. For our multi-relational clinical data, we utilise Relational Graph Convolutional Networks (RGCN)³⁹, designed for highly multi-relational data, to model patient clinical KG embeddings.

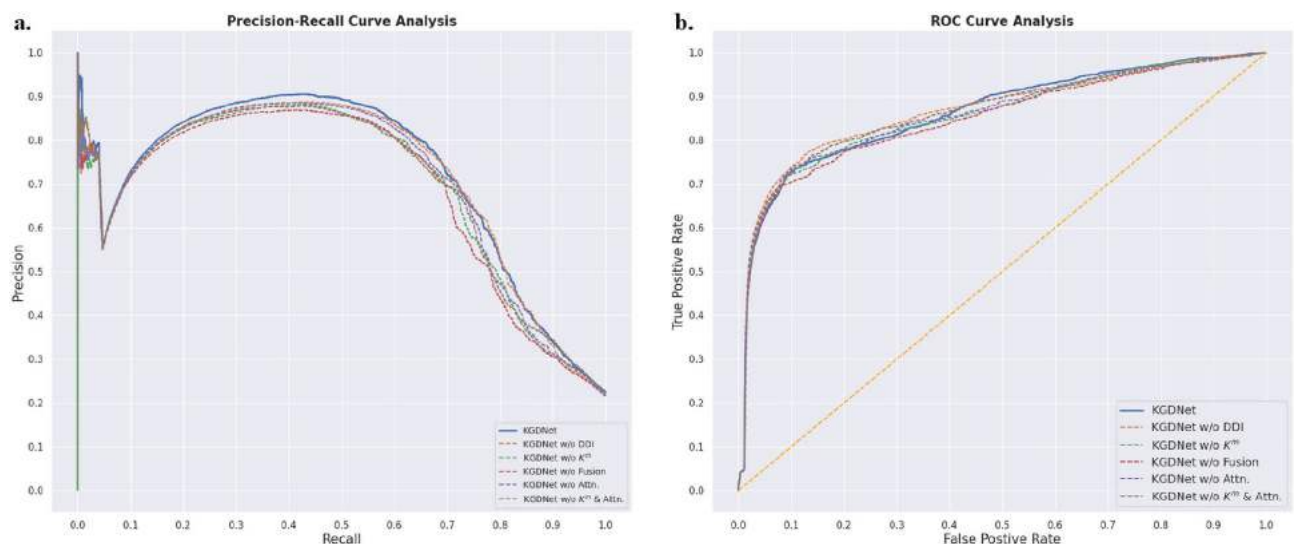


Fig. 4. Comparison of (a) Precision-Recall Curves and (b) ROC-AUC curves of KGDNet along with various ablated variants. KGDNet is represented by the solid blue plot-line.

Furthermore, we propose a refined approach of developing medication-specific KGs for each patient admission. Instead of creating a shared “memory bank” for all patients, as used in^{3–7}, we create individual KGs containing known drug co-occurrences associated with prescribed medications. This strategy aims to enhance the efficacy of recommendation sets, aligning more closely with the personalised nature of healthcare interventions. However, we define a global Drug-Drug Interaction (DDI) KG, similar to GAMENet³, comprising of medicines, with edges denoting the DDIs. Employing GNNs, we learn this KG’s medical representations. These representations are integrated with the medication features of an admission, to derive the unified medication features for that specific admission. Subsequently, these combined features are employed in RNNs to effectively capture and model the patient’s medication history. We use Graph Convolutional Networks (GCN) to learn the medication KGs and the DDI KG. After obtaining the embeddings through GNN, they are then passed onto RNNs to learn the temporal data of the patient across admissions.

Gated recurrent units for longitudinal data

From existing works we have identified that considering historical data along with current clinical information is important^{3,14}. While GAMENet³ and SafeDrug⁴ use RNNs To learn historical data, COGNet⁶ and ACDNet⁷ make use of Self-Attention mechanisms. We make use of Recurrent Neural Networks (RNNs), a bidirectional neural network that can effectively process sequential data. Gated Recurrent Units (GRUs)⁴⁰ are a commonly used RNN that uses two gates - “reset” gate and “update” gate - that determine whether to pass features to the output or forget them. Instead of using RNNs directly on the multi-hot vector clinical data as done by^{3,4}, we first create KGs, then apply GNNs and finally apply GRU on the KG embeddings.

In order to capture the temporal features of each visit, we need to comprehensively gain information about both the clinical and medication history of the patient. Previous works have used simple concatenation techniques that combine the clinical RNN features with the KG embeddings from the global medicine memory bank. We utilise a Fusion layer that performs collaborative filtering⁴¹ by taking the clinical and medication RNN streams for an admission obtained from the GRU cells and converts them to a joint medical stream that effectively captures the temporal features of the patient’s diagnoses, procedure and prescription data. We concatenate the two streams and pass the result through a series of convolutions and a Multi-Layer Perceptron (MLP) to obtain the joint medical features, which are passed into another RNN for learning the joint temporal stream.

Attention mechanism for generating recommendations

After obtaining the joint features for a patient, like COGNet⁶, we then use an Attention mechanism for our final recommender module to effectively learn from the temporal features extracted from the fusion module. An attention mechanism is a neural network architecture that computes the relevance of each element in a sequence to every other element, allowing the model to focus on relevant portions of the input by assigning weights. Attention mechanisms have been frequently used with GNNs^{6,7,42,43} in healthcare field. However, previous works like COGNet⁶ and ACDNet⁷ have used Transformer Encoders, i.e., self-attention followed by layer normalisation¹⁶, to learn the patient’s clinical history. We propose using Multi-Head Attention (MHA)¹⁵ instead of self-attention based Transformers, enabling us to provide various inputs related to different medical states of the patient to the attention module. MHA take in query, key and value objects. The joint medical stream is assigned to key and value while the current clinical temporal features are assigned to the query object to put emphasis on the current clinical state along with the patient history. The result obtained from MHA is then passed onto a layer normalisation module along with the clinical embeddings of the current state to provide emphasis on the current admission. The result of the layer normalisation is then passed to a MLP followed by a Tanh activation function to return a set of weights upon which a threshold is applied to generate the recommendation set.

LLMs in medical recommendation systems

We also explore the use of Large Language Models (LLMs) in medical recommendation systems to model longitudinal patient data in comparison with our approach of using GNNs. The strong performance of LLMs like GPT⁴⁴ in recent times has led them to be utilised in hugely diverse scenarios including the healthcare field⁴⁵ and for recommendation systems⁴⁶. Hence it becomes crucial that we look into the advantages and drawbacks of LLMs in the medicine recommendations. However, it has been found that GPTs and LLMs still have considerable room for improvement in areas such as information gathering and adhering to guidelines⁴⁷. GraphCare⁴⁸ utilises LLMs and external biomedical KGs to build patient-specific KGs, which are then used to train our proposed Bi-attention Augmented GNNs for drug recommendation. LEADER⁴⁹ uses LLMs by employing custom prompts and a novel output layer. It transfers LLM knowledge to a smaller distilled model, balancing power with efficiency. Both models give good results on MIMIC-IV. However, a major drawback is that both models do not take into account DDIs between drugs in the recommendation sets. Both models also have shortcomings related to hallucinations, biases⁴⁸ and high computational costs⁴⁹ that could compromise the effectiveness of these models.

Measurement and performance of KGDNet

We employ various metrics such as Precision-Recall Area Under Curve (PRAUC), F1 score, and Jaccard score are utilised to evaluate the predictive performance of KGDNet against multiple baselines. A custom DDI rate metric is introduced to evaluate the KGDNet’s effectiveness in controlling DDIs within the recommendation set, drawing inspiration from the methodology introduced by SafeDrug⁴. We demonstrate that KGDNet outperforms the baselines in every performance metric on the MIMIC-IV cohort. We also evaluate the performance of our model over a range of DDI thresholds to demonstrate that our model is capable of controlling the DDI rate while also maintaining a high level of performance. We conduct an ablation study to analyse the impact of various individual modules within KGDNet. Finally, we conduct a case study by selecting a patient from the test dataset

and evaluate our model and the baselines’ approaches against the ground truth, i.e, the patient’s prescription. We show that KGDNet is successful in providing reliable recommendations reflecting the ground truth in real-world healthcare situations.

Methods

Dataset description

We consider EHR data from a benchmark inpatient dataset, MIMIC-IV⁹, along with medical ontology data from ICD and DDI knowledge from TWOSIDES¹¹. Table 4 provides some dataset statistics while details of the dataset and preprocessing can be found in 1.1. in Supplementary Information (Fig. 5).

Problem formulation

In our medicine recommendation system, we will utilise Electronic Health Records of patients as well as external medical information like Drug-Drug Interaction Data, ontologies, etc. to create a deep learning model that will provide personalized medicine recommendations to patients.

Electronic health records

Electronic Health Records (EHRs) are used to store medical histories of patients. This is done using longitudinal vectors that store information about a patient’s diagnoses, procedures and drugs prescribed to them.

For a patient n , the EHR can be represented as $X^n = \{X_n^1, \dots, X_n^t, \dots, X_n^T\}$, with T referring to the total number of admissions. For each t -th admission of the i -th patient, $X_i^t = \{d^t, p^t, m^t\}$ consists of vectors for diagnosis $d^t \in |\mathcal{D}|$, procedure $p^t \in |\mathcal{P}|$ and medicine codes $m^t \in |\mathcal{M}|$.

As the diagnoses and procedures in the EHR are uniquely defined in the ICD ontology, we can integrate the diagnoses and procedure sets into one combined set we will call as clinical set, \mathcal{C} , such that $c^t \in |\mathcal{C}|$, where $\mathcal{C} = \mathcal{D} \cup \mathcal{P}$. This will help us establish a variety of relations between the various diagnoses and procedures that a patient is or has been associated with on a patient and cohort level.

Patient medical knowledge graphs

For each admission we then generate two disjoint knowledge graphs, K_c^t for clinical information and K_m^t for medication information. For the clinical KG, we incorporate various relations such as patient diagnoses and patient procedures along with diagnoses and procedures related to them to capture extensive information about the patient’s conditions by encoding information under the ICD ontology. For the medication KG, we incorporate relations such as patient prescriptions along with related medicines under the ATC ontology.

We then transform each medicine, procedure and diagnoses nodes using embeddings to acquire node features.

$$V_c^t = E_c \cdot c^t \tag{1}$$

$$V_m^t = E_m \cdot m^t \tag{2}$$

where $E_{\{c,m\}} \in \mathbb{R}^{n \times |\mathcal{C}|, |\mathcal{M}|}$ and n is the embedding size. We augment the set of nodes in each graph by one to denote the patient node. Fig. 6a. shows a visualisation of a sample Knowledge Graph of a patient’s admission. The various diagnoses, procedures and the different relations between them are depicted.

DDI knowledge graph

We also create a DDI knowledge graph, K_{ddi} consisting of medical nodes and DDI relations between the nodes. This graph captures all the possible DDI pairs in our medicine dataset. The medical nodes, also encoded using the ATC ontology, belong to the same medicine set \mathcal{C} . The nodes features are extracted using the same embedding methods used in K_m^* . The DDI KG we used is visualised in Fig. 6b.

Items	Size
# of Patients	75535
# of Admissions	194883
# of Diagnoses	2007
# of Procedures	1500
# of Medicines	146
Avg./Max # of Admissions	2.45/66
Avg./Max # of Diagnoses	6.45/228
Avg./Max # of Procedures	2.24/72
Avg./Max # of Medicines	9.12/72
# of DDI pairs	519

Table 4. Dataset statistics.

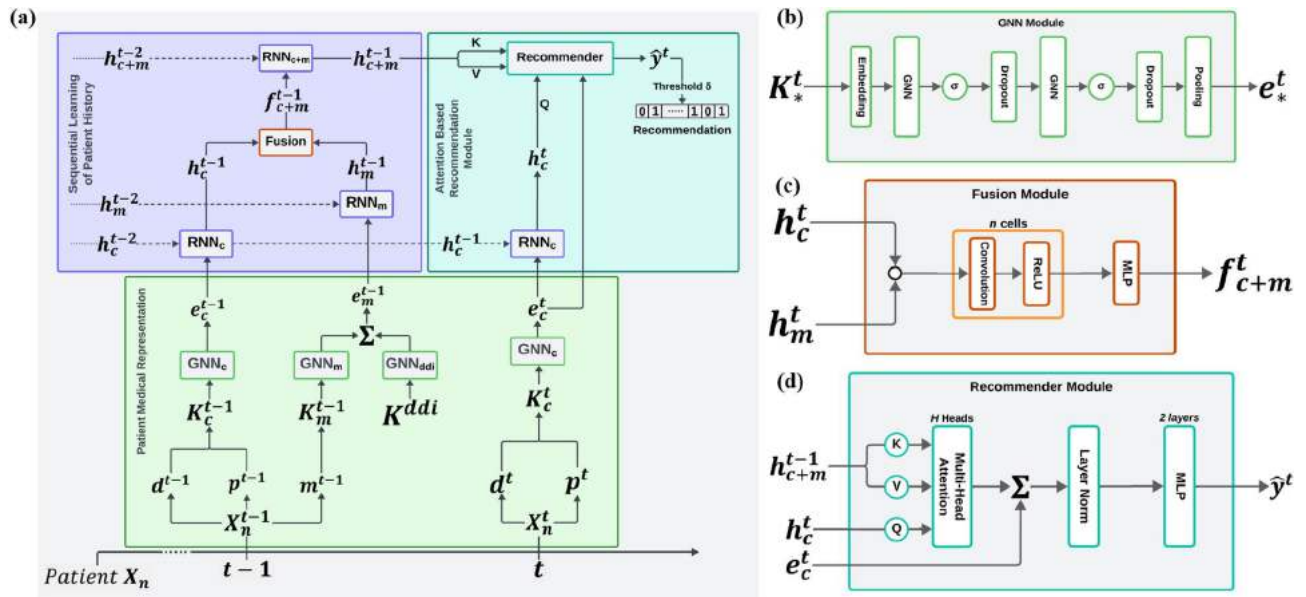


Fig. 5. KGDNet Framework. (a) In the patient medical representation phase, we create medical KGs for each admission. Using diagnoses and procedure data we generate clinical KGs and using medication data, we create medicine KGs. Embeddings from medicine KGs are subtracted from DDI KG embeddings. In sequential learning of patient history, we learn the hidden temporal features of each admission using RNNs. We generate the hidden features for the clinical and medicine streams. We generate joint hidden medical streams from the clinical and medicine streams using a Fusion Module that generates joint features which are passed onto a joint RNN. In the recommendation phase, we use an Attention-based recommender module that utilises MHA which takes the joint hidden features of the previous admission along with current hidden clinical features. The output from MHA is added with the current clinical embeddings and normalised to get the recommendations. (b) Graph Neural Network model used in our framework. For clinical KGs, the GNN used is R-GCN and for medicine and DDI KGs, the GCN is used. σ in the image signifies the ReLU activation function. (c) Fusion module for fusing the clinical and medicine RNN streams to generate joint medical features. The circle signifies that the clinical and medical hidden features for the admission t are concatenated. (d) The recommender module for prescribing the medication set. The clinical embedding for the current admission is added to the output from MHA and then passed to the Layer Normalisation layer.

Medication recommendation

Given the healthcare record of patients in the form of most recent clinical KG K_c^t along with their past medical KGs $K_c^{1:t-1}$ and $K_m^{1:t-1}$ and DDI KG K^{ddi} , our proposed model aims to recommend a set of medicines by generating a multi-class output $\hat{y}^t \in \{0, 1\}^{|\mathcal{M}|}$ of medicines while minimising the Drug-Drug Interactions between them.

Patient representation using knowledge graph embedding

In order to learn our medical KGs representations, we resort to variants of Graph Neural Networks. Several GNNs such as RGCN³⁹, GCN²¹ have been designed specifically for heterogeneous data and can efficiently account for various relations within a graph. We designate Relational Graph Convolutional Network (R-GCN) to model multi-relational data in learning the node embeddings for clinical KGs. R-GCN is used to learn the node embeddings as it enables us to apply Graph Convolutional Networks on data that has a number of relations to be accounted for. Node-wise formulation for R-GCN works as below:

$$V_c^{i,t} = \Theta_{\text{root}} \cdot V_i + \sum_{r \in \mathcal{R}} \sum_{j \in \mathcal{N}_r(i)} \frac{1}{|\mathcal{N}_r(i)|} \Theta_r \cdot V_c^j \quad (3)$$

where V_c^i denotes the embeddings for clinical node i and \mathcal{R} denotes the set of relations, i.e., edge types. Edge type needs to be a one-dimensional vector which stores a relation identifier $\in \{0, \dots, |\mathcal{R}| - 1\}$ for each edge.

We designate Graph Convolutional Network (GCN) to model learn the embeddings of medication KGs and the DDI KG. GCN efficiently learns embedding of KGs with few relational types and enables us to assign weights to different types of edges, thus enabling us to assign higher weights to edges corresponding to prescriptions and assign lower weights to DDI edges in order to avoid them during recommendation. Its node-wise formulation is given by:

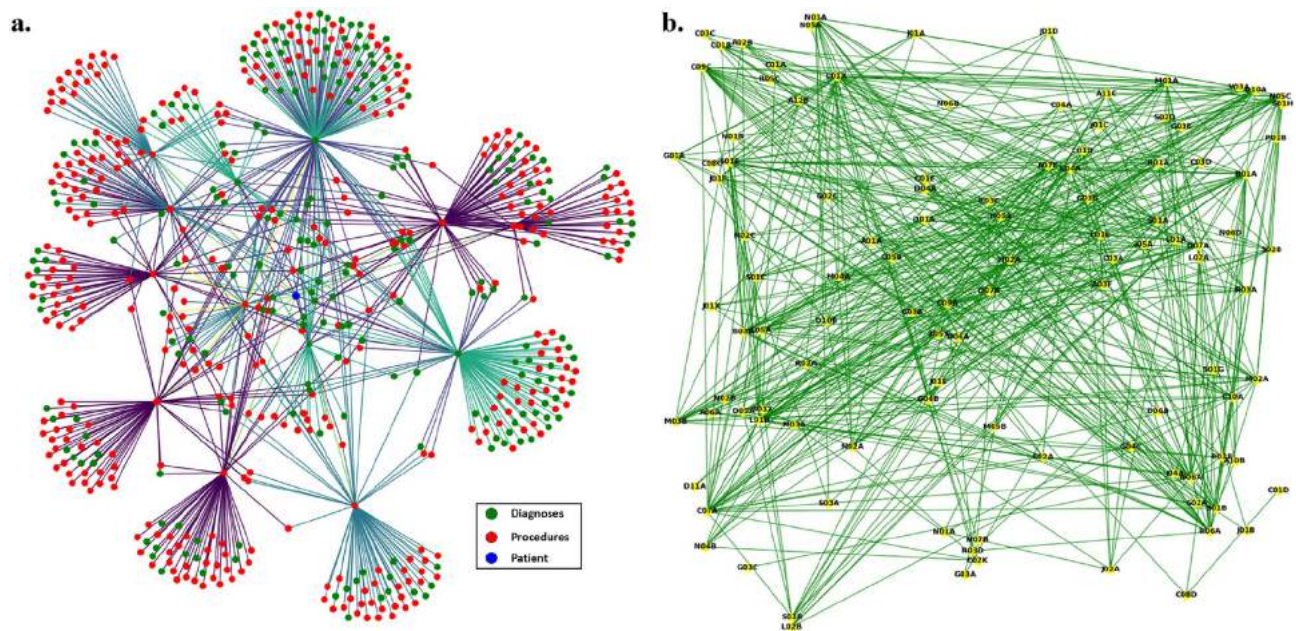


Fig. 6. Visualisation of Knowledge Graphs. **(a)** Clinical knowledge graph of a patient's admission record. The nodes represent different diagnoses and procedures along with an auxiliary patient node. The various relations between the nodes are represented in different colours. **(b)** The global DDI knowledge graph used in KGDNet. The nodes are the medicines from our cohort, represented using ATC-3 ontology, which have DDIs with other medicines. The relations in both graph are bidirectional which enables backpropagation of messages in GNNs during training.

$$\mathbf{v}_m^{i,t} = \Theta^\top \sum_{j \in \mathcal{N}(i) \cup \{i\}} \frac{e_{j,i}}{\sqrt{\hat{d}_j \hat{d}_i}} \mathbf{v}_m^j \quad (4)$$

with $\hat{d}_i = 1 + \sum_{j \in \mathcal{N}(i)} e_{j,i}$, where $e_{j,i}$ denotes the edge weight from source node j to target node i . To learn the medication node embeddings of a patient's admission, we perform weighted sum on the admission's medication node embeddings and the DDI KG embeddings to fuse the two KGs together in order to account for both the patient's prescriptions and the DDIs that might occur between them.

$$\hat{\mathbf{V}}_m^t = \mathbf{V}_m^t - \mathbf{V}_{ddi} \quad (5)$$

where \mathbf{V}_m^t refers to the node embeddings for admission t , and \mathbf{V}_{ddi} refers to the node embeddings of the DDI KG. For each type of graph, as shown in Fig. 1, we apply two layers of GNNs to learn the node embeddings for each clinical and medication knowledge graph. The embeddings are then aggregated through a readout mechanism to obtain a tuple of medical representations of a patient's admission embeddings $\{e_c^{1:t}, e_m^{1:t-1}\}$, for clinical and medication KG embeddings, respectively, that are then forwarded to the sequential learning mechanism. The architecture of the GNN module is illustrated in Fig. 5b.

Sequential learning of patient history

Given the tuple of KG embeddings for each patient, we learn the temporal features from the patient's admissions using Recurrent Neural Networks (RNNs) on both the clinical and medication embeddings. We then fuse the hidden features obtained from both RNNs using a fusion module to get the combined hidden features of the clinical and medicine streams to learn the joint temporal features using another RNN model for combined features.

Learning patient history using recurrent neural networks

Using the medical embeddings obtained through GNNs, we learn the hidden features of each admission by utilising two separate RNNs, RNN_c and RNN_m to encode admission-wise clinical and medicine data as follows:

$$h_*^t = RNN_*(e_*^t, h_*^{t-1}) \quad (6)$$

For our RNNs, we use Gated Recurrent Unit (GRU)⁴⁰ cells for learning the hidden features, following the example of SafeDrug⁴ and GAMENet³.

Fusion of clinical and medication history

After generating hidden temporal features for each admission prior to the current, t , we then use a feed-forward fusion mechanism to combine the clinical and medicine streams. This is achieved using a two-step process. The first involves concatenating the clinical and medical streams and then passing them through n series of convolution functions and then fed into a MLP layer. The MLP is used to combine the information from the medicine and clinical stream to obtain a more comprehensive and informative feature representation. It consists of two layers of projection matrices, P_1 and P_2 with a ReLU activation function between them to obtain the fused features f_{c+m} .

$$f_{c+m}^t = P_1(P_2(h_c^t \cdot h_m^t)) \quad (7)$$

The fused features of each admission before the current one are passed through an RNN, RNN_{c+m} , to learn the joint hidden features of the combined streams. The architecture of the Fusion module is illustrated in Fig. 5c.

$$h_{c+m}^t = RNN_{c+m}(f_{c+m}^t, h_{c+m}^{t-1}) \quad (8)$$

Attention mechanism for generating recommendations

Attention module

After completing the above steps, we have now obtained h_c^t , representing the patient's hidden clinical features and h_{c+m}^{t-1} , representing joint hidden features, for each admission record t . The above features capture important information about a patient's history and current information. These features can be used to provide accurate, personalised recommendations. We use Multi-Head Attention (MHA), to learn the clinical and joint hidden features. The architecture of the attention-based recommendation module is illustrated in Fig. 5d.

$$MHA(q, k, v; \mathcal{H}) = W_{MHA}[head_1, head_2, \dots, head_{\mathcal{H}}] \quad (9)$$

$$head_i = \text{attention}(W_i^q q, W_i^k k, W_i^v v) \quad (10)$$

where query q is h_c^t , key k and value v are h_{c+m}^{t-1} and \mathcal{H} is the number of heads.

Decoder

Furthermore, we intend to leverage the clinical embeddings, e_c^t , associated with the present admission, as demonstrated by GAMENet³. This approach ensures a comprehensive integration of the current clinical information, allowing us to factor in the current clinical state and address any potential new illnesses that may not have been documented in prior admissions. We do this by adding the current clinical embeddings with the result of MHA on the hidden features and then performing layer normalisation on the result, taking inspiration from COGNet⁶.

$$l^t = \text{LayerNorm}(e_c^t + MHA(h_c^t, h_{c+m}^{t-1}, h_{c+m}^{t-1})) \quad (11)$$

Recommendation

Finally, we pass the result to an MLP module, consisting of two projection layers P_m^1, P_m^2 , to convert the result of normalisation, l^t , to a set of scores corresponding to the medicine set. This set of weights is then passed through a Tanh-activation layer to return the set of recommended medicines, \hat{y}^t .

$$\hat{y}^t = \text{Tanh}(P_m^2(\sigma(P_m^1(l^t)))) \quad (12)$$

Training and baselines

Objective functions

The recommendation task has been designed as a multi-label binary classification problem. With the size of the medication set as $|\mathcal{M}|$, m^t denotes the ground truth medication set at the t -th visit and \hat{m}^t denotes the set of recommended medicines.

Require: Training Patient Records $X = \{X_1, \dots, X_N\}$, DDI Knowledge Graph K^{ddi} , DDI Adjacency matrix A_{ddi} , DDI threshold λ

- 1: **for** patient $i = 1$ to N **do**
- 2: Read the patient i 's admission records $X_i^1, \dots, X_i^t, \dots, X_i^T$
- 3: **for** admission $t = 1$ to T **do**
- 4: Read diagnoses d^t , procedures p^t and medication m^t for t -th admission.
- 5: Generate KG K_c^t using patient's clinical information and relations between diagnoses and procedures.
- 6: Generate KG K_m^t using patient's medication and drug co-occurrence information.
- 7: Generate clinical and medication KG embeddings $\{e_c^t, e_m^t\}$ using GCNs.
- 8: Store the KG embeddings in patient's medical historical features for temporal learning.
- 9: Generate clinical and medicine temporal hidden features $\{h_c^t, h_m^{t-1}\}$ using RNNs in Equation (6).
- 10: Generate joint historical medical stream h_{c+m}^{t-1} for admission t by fusing historical clinical and medicine longitudinal streams in Equations (7), (8).
- 11: Generate normalised output l^t by combining current clinical embeddings e_c^t with Multi-Head Attention on joint hidden features h_{c+m}^{t-1} and current clinical features h_c^t in Equation (9), (11).
- 12: Generate recommendation output by passing l^t through a MLP and TanH activation function in Equation (12).
- 13: **end for**
- 14: Apply threshold to generate the recommendation set.
- 15: Calculate losses L_{bce} (13), L_{multi} (14), L_{ddi} (15) using recommendation set.
- 16: **end for**
- 17: Calculate combined losses L_{total} (16) and update model parameters.

Algorithm 1. One training epoch of KGDNet

Prediction Loss Functions: We use binary cross entropy loss and multi-label margin loss to evaluate the prediction loss:

$$L_{bce} = - \sum_{t=1}^T \sum_{i=1}^{|\mathcal{M}|} y_i^t \cdot \log(\hat{y}_i^t) + (1 - y_i^t) \cdot (1 - \log(\hat{y}_i^t)) \quad (13)$$

$$L_{multi} = \sum_{t=1}^T \sum_{i,j=1}^{|\mathcal{M}|} \frac{\max(0, 1 - (\hat{y}_i^t[y_i^t] - \hat{y}_j^t[i]))}{|\mathcal{M}|} \quad (14)$$

DDI Loss Function: We also introduce a DDI loss to ensure safety while recommending medicines by controlling the amount of DDI present in a set of recommendations. To calculate DDI loss, we introduce DDI adjacency matrix A_{ddi} .

$$L_{ddi} = \sum_{t=1}^T \sum_{i,j=1}^{|\mathcal{M}|} A_{ddi} \cdot \hat{y}_i^t \cdot y_j^t \quad (15)$$

Combined Loss Function: We use weighted sum of the multiple loss functions as introduced by Dosovitskiy and Djolonga⁵⁰ when training deep learning models. We utilise the approach introduced by SafeDrug⁴.

$$L_{total} = \Phi_{ddi}(\Phi_{pred}L_{bce} + (1 - \Phi_{pred})L_{multi}) + (1 - \Phi_{ddi})L_{ddi} \quad (16)$$

Where Φ_{pred} and Φ_{ddi} are pre-defined hyperparameters. Φ_{pred} weighs the prediction losses and is selected arbitrarily. Φ_{ddi} is determined during the training process by observing the patient-level DDI rate, X_{ddi} , for a patient X . If X_{ddi} is below a threshold λ , then we will adjust our loss function to focus only on maximising prediction accuracy, otherwise λ will adjust our loss function to reduce DDI. Φ_{ddi} is determined by the function,

$$\Phi_{ddi} = \begin{cases} 1, & X_{ddi} < \lambda \\ \max(0, 1 - \frac{X_{ddi} - \lambda}{\rho}), & X_{ddi} \geq \lambda \end{cases} \quad (17)$$

where ρ is the correcting factor. We can change Λ to change the DDI rate in our recommendations.

Baselines

We evaluate our model against various EHR based medication recommendation models. These baselines can be categorised into two main categories, *instance-based*: LR and LEAP¹, that do not utilise past medical data of patients and *longitudinal-based*^{3–7,14,17–19,22,23} that utilise the historical data. Our model is compared to each of the baselines using PRAUC, Jaccard, F1 scores, DDI Rates and average number of medicines prescribed as metrics.

- **LR** is Logistic Regression with L2 regularisation
- **LEAP**¹ is an instance-based sequential decision-making recommendation model
- **RETAIN**¹⁴ uses a two-level attention model to identify important past visits and clinical variables.
- **GAMENet**³ leverages graph-augmented memory networks by applying a fusion-based GCN on drug co-occurrences and DDIs and then performs attention-based memory search using queries from patient records.
- **SafeDrug**⁴ performs medication recommendation by considering DDIs between medicines and their molecular structure using a “message passing neural network”.
- **MICRON**⁵ uses a “recurrent residual network” that sequentially updates and propagates a patient’s medical information.
- **COGNet**⁶ utilises the Transformer Encoder mechanism to learn the patient’s history and then uses a copy-or-predict mechanism that balances historical records along with current clinical information to provide recommendations.
- **4SDrug**¹⁷ uses set-oriented representations, similarity metrics, and importance-based set aggregation for symptom analysis, along with intersection-based set augmentation for recommending medications.
- **DrugRec**¹⁸ uses causal inference to address bias, track patients’ health history, and manage DDIs by solving them as a satisfiability problem.
- **MoleRec**¹⁹ employs a hierarchical architecture to model molecular inter-substructure interactions and their on a patient’s health condition.
- **ACDNet**⁷ uses the Transformer Encoder¹⁵ mechanism to learn the patient’s clinical and medication history along with GCN²¹ to learn the global medicine data and uses cosine similarity to provide recommendations.
- **PROMISE**²² encodes EHR hypergraphs to learn patient representations and acquires patient representations with semantic information by encoding clinical texts from EHR.
- **VITA**²³ introduces two novel ideas: relevant visit selection and target aware attention to capture the relevancy between current and past visits.

Data availability

The EHR cohort used in this paper, MIMIC IV v2.2, is publicly available to credentialed users after signing the data use agreement on the website. <https://physionet.org/content/mimiciv/2.2/>

Code availability

The original KGDNet code and other codes used in this work are publicly available on <https://github.com/Rajat1206/KGDNet>.

Received: 14 May 2024; Accepted: 8 October 2024

Published online: 26 October 2024

References

1. Zhang, Y., Chen, R., Tang, J., Stewart, W.F., & Sun, J. Leap: Learning to prescribe effective and safe treatment combinations for multimorbidity. In: *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining. KDD '17*, pp. 1315–1324. Association for Computing Machinery, New York, NY, USA (2017). <https://doi.org/10.1145/3097983.3098109>
2. Gong, F., Wang, M., Wang, H., Wang, S. & Liu, M. Smr: Medical knowledge graph embedding for safe medicine recommendation. *Big Data Res.* **23**, 100174. <https://doi.org/10.1016/j.bdr.2020.100174> (2021).
3. Shang, J., Xiao, C., Ma, T., Li, H., & Sun, J. GAMENet: Graph Augmented MEmory Networks for Recommending Medication Combination. arXiv preprint [arXiv:1809.01852](https://arxiv.org/abs/1809.01852) (2018).
4. Yang, C., Xiao, C., Ma, F., Glass, L., & Sun, J. SafeDrug: Dual Molecular Graph Encoders for Recommending Effective and Safe Drug Combinations (2022).
5. Yang, C., Xiao, C., Glass, L. & Sun, J. Change Matters: Medication Change Prediction with Recurrent Residual Networks (2021)
6. Wu, R., Qiu, Z., Jiang, J., Qi, G. & Wu, X. Conditional generation net for medication recommendation. In: *Proceedings of the ACM Web Conference 2022. WWW '22*, pp. 935–945. Association for Computing Machinery, New York, NY, USA (2022). <https://doi.org/10.1145/3485447.3511936>.
7. Mi, J., Zu, Y., Wang, Z. & He, J. Acdnet: Attention-guided collaborative decision network for effective medication recommendation. *J. Biomed. Inf.* **149**, 104570. <https://doi.org/10.1016/j.jbi.2023.104570> (2024).
8. Johnson, A. E. et al. MIMIC-III, a freely accessible critical care database. *Sci. Data* **3**(1), 1–9 (2016).
9. Johnson, A. E. et al. MIMIC-IV, a freely accessible electronic health record dataset. *Sci. Data* **10**(1), 1 (2023).
10. Huang, X., Zhang, J., Xu, Z., Ou, L. & Tong, J. A knowledge graph based question answering method for medical domain. *PeerJ Comput. Sci.* **7**, 667 (2021).
11. Tatonetti, N. P., Ye, P. P., Daneshjou, R. & Altman, R. B. Data-driven prediction of drug effects and interactions. *Sci. Trans. Med.* **4**(125), 125–3112531. <https://doi.org/10.1126/scitranslmed.3003377> (2012).
12. Smithburger, P. L., Kane-Gill, S. L. & Seybert, A. L. Drug-drug interactions in the medical intensive care unit: an assessment of frequency, severity and the medications involved. *Int. J. Pharmacy Practice* **20**(6), 402–408 (2012).
13. Nyamabo, A. K., Yu, H. & Shi, J.-Y. SSI-DDI: substructure-substructure interactions for drug-drug interaction prediction. *Briefings Bioinf.* **22**(6), 133. <https://doi.org/10.1093/bib/bbab133> (2021).

14. Choi, E., Bahadori, M.T., Sun, J., Kulas, J., Schuetz, A. & Stewart, W. Retain: An interpretable predictive model for healthcare using reverse time attention mechanism. *Adv. Neural Inf. Process. Syst.* **29** (2016)
15. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, L. & Polosukhin, I. Attention Is All You Need (2023)
16. Ba, J.L., Kiros, J.R. & Hinton, G.E. Layer Normalization (2016)
17. Tan, Y., Kong, C., Yu, L., Li, P., Chen, C., Zheng, X., Hertzberg, V.S. & Yang, C. 4sdrug: Symptom-based set-to-set small and safe drug recommendation. In: *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining. KDD '22*, pp. 3970–3980. Association for Computing Machinery, New York, NY, USA (2022). <https://doi.org/10.1145/3534678.3539089>
18. Sun, H. et al. Debiased, longitudinal and coordinated drug recommendation through multi-visit clinic records. *Adv. Neural Inf. Process. Syst.* **35**, 27837–27849 (2022).
19. Yang, N., Zeng, K., Wu, Q. & Yan, J. Molerec: Combinatorial drug recommendation with substructure-aware molecular representation learning. In: *Proceedings of the ACM Web Conference 2023. WWW '23*, pp. 4075–4085. Association for Computing Machinery, New York, NY, USA (2023). <https://doi.org/10.1145/3543507.3583872>
20. Veličković, P., Cucurull, G., Casanova, A., Romero, A., Liò, P. & Bengio, Y. Graph Attention Networks (2018)
21. Kipf, T.N. & Welling, M. Semi-Supervised Classification with Graph Convolutional Networks (2017)
22. Wu, J., Yu, X., He, K., Gao, Z. & Gong, T. Promise: A pre-trained knowledge-infused multimodal representation learning framework for medication recommendation. *Inf. Process. Manage.* **61**(4), 103758 (2024).
23. Kim, T., Heo, J., Kim, H., Shin, K. & Kim, S.-W. Vita: 'carefully chosen and weighted less' is better in medication recommendation. In: *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 38, pp. 8600–8607 (2024)
24. Abul-Husn, N. S. & Kenny, E. E. Personalized medicine and the power of electronic health records. *Cell* **177**(1), 58–69 (2019).
25. Menachemi, N. & Collum, T.H. Benefits and drawbacks of electronic health record systems. *Risk management and healthcare policy*, 47–55 (2011)
26. Nigo, M. et al. Deep learning model for personalized prediction of positive mrsa culture using time-series electronic health records. *Nat. Commun.* **15**(1), 2036 (2024).
27. Li, L. et al. Real-world data medical knowledge graph: construction and applications. *Artif. Intell. Med.* **103**, 101817. <https://doi.org/10.1016/j.artmed.2020.101817> (2020).
28. Lee, E., Lee, D., Baek, J. H., Kim, S. Y. & Park, W.-Y. Transdiagnostic clustering and network analysis for questionnaire-based symptom profiling and drug recommendation in the uk biobank and a korean cohort. *Sci. Rep.* **14**(1), 4500 (2024).
29. Zhang, Y. et al. Emerging drug interaction prediction enabled by a flow-based graph neural network with biomedical network. *Nat. Comput. Sci.* **3**(12), 1023–1033 (2023).
30. Wang, Y., Yang, Z. & Yao, Q. Accurate and interpretable drug-drug interaction prediction enabled by knowledge subgraph learning. *Commun. Med.* **4**(1), 59 (2024).
31. Wishart, D. S. et al. DrugBank: A comprehensive resource for in silico drug discovery and exploration. *Nucleic Acids Res.* **34**, 668–672. <https://doi.org/10.1093/nar/gkj067> (2006).
32. Dai, Y., Wang, S., Xiong, N. N. & Guo, W. A survey on knowledge graph embedding: Approaches, applications and benchmarks. *Electronics* [SPACE] <https://doi.org/10.3390/electronics9050750> (2020).
33. Corso, G., Stark, H., Jegelka, S., Jaakkola, T. & Barzilay, R. Graph neural networks. *Nat. Rev. Methods Primers* **4**(1), 17 (2024).
34. Zhou, J. et al. Graph neural networks: A review of methods and applications. *AI Open* **1**, 57–81 (2020).
35. He, Q., Li, X. & Cai, B. Graph neural network recommendation algorithm based on improved dual tower model. *Sci. Rep.* **14**(1), 3853 (2024).
36. Gilmer, J., Schoenholz, S.S., Riley, P.F., Vinyals, O. & Dahl, G.E. Neural Message Passing for Quantum Chemistry (2017)
37. Shang, J., Ma, T., Xiao, C. & Sun, J. Pre-training of graph augmented transformers for medication recommendation. arXiv preprint [arXiv:1906.00346](https://arxiv.org/abs/1906.00346) (2019)
38. Liu, T., Shen, H., Chang, L., Li, L. & Li, J. Iterative heterogeneous graph learning for knowledge graph-based recommendation. *Sci. Rep.* **13**(1), 6987 (2023).
39. Schlichtkrull, M. et al. Modeling relational data with graph convolutional networks. In *The Semantic Web* (eds Gangemi, A. et al.) 593–607 (Springer, 2018).
40. Cho, K., Merriënboer, B., Gulcehre, C., Bahdanau, D., Bougares, F., Schwenk, H. & Bengio, Y. Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation (2014)
41. He, X., Liao, L., Zhang, H., Nie, L., Hu, X. & Chua, T.-S. Neural collaborative filtering. In: *Proceedings of the 26th International Conference on World Wide Web. WWW '17*, pp. 173–182. International World Wide Web Conferences Steering Committee, Republic and Canton of Geneva, CHE (2017). <https://doi.org/10.1145/3038912.3052569>
42. Wang, S., Qiao, J. & Feng, S. Prediction of lncrna and disease associations based on residual graph convolutional networks with attention mechanism. *Sci. Rep.* **14**(1), 5185 (2024).
43. Hasibi, R., Michael, T. & Oyarzún, D. A. Integration of graph neural networks and genome-scale metabolic models for predicting gene essentiality. *npj Syst. Biol. Appl.* **10**(1), 24 (2024).
44. Radford, A. Improving language understanding by generative pre-training (2018)
45. Bedi, S., Jain, S.S. & Shah, N.H. Evaluating the clinical benefits of llms. *Nature Medicine*, 1–2 (2024)
46. Wu, L. et al. A survey on large language models for recommendation. *World Wide Web* **27**(5), 60 (2024).
47. Hager, P., Jungmann, F., Holland, R., Bhagat, K., Hubrecht, I., Knauer, M., Vielhauer, J., Makowski, M., Braren, R. & Kaissis, G., et al. Evaluation and mitigation of the limitations of large language models in clinical decision-making. *Nature medicine*, 1–10 (2024)
48. Jiang, P., Xiao, C., Cross, A. & Sun, J. GraphCare: Enhancing Healthcare Predictions with Personalized Knowledge Graphs (2024). [arxiv:2305.12788](https://arxiv.org/abs/2305.12788)
49. Liu, Q., Wu, X., Zhao, X., Zhu, Y., Zhang, Z., Tian, F. & Zheng, Y. Large Language Model Distilling Medication Recommendation Model (2024). [arxiv:2402.02803](https://arxiv.org/abs/2402.02803)
50. Dosovitskiy, A. & Djolonga, J. You only train once: Loss-conditional training of deep networks. In: *International Conference on Learning Representations* (2019)

Author contributions

S.S. conceptualized and proposed the project. R.M. developed the KGDNet model, including designing its architecture and implementing the code. R.M. also conducted model training, executed the experiments, and performed comprehensive data analysis. The initial draft of the manuscript was prepared by R.M. Both S.S. and R.M. critically reviewed, revised, and edited the manuscript to its final form. S.S. provided overall project supervision and guidance throughout the research process.

Funding

Not Applicable

Open access funding provided by Vellore Institute of Technology.

Declarations

Competing interests

The authors declare no competing interests.

Additional information

Supplementary Information The online version contains supplementary material available at <https://doi.org/10.1038/s41598-024-75784-5>.

Correspondence and requests for materials should be addressed to S.S.

Reprints and permissions information is available at www.nature.com/reprints.

Publisher's note Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.

Open Access This article is licensed under a Creative Commons Attribution 4.0 International License, which permits use, sharing, adaptation, distribution and reproduction in any medium or format, as long as you give appropriate credit to the original author(s) and the source, provide a link to the Creative Commons licence, and indicate if changes were made. The images or other third party material in this article are included in the article's Creative Commons licence, unless indicated otherwise in a credit line to the material. If material is not included in the article's Creative Commons licence and your intended use is not permitted by statutory regulation or exceeds the permitted use, you will need to obtain permission directly from the copyright holder. To view a copy of this licence, visit <http://creativecommons.org/licenses/by/4.0/>.

© The Author(s) 2024